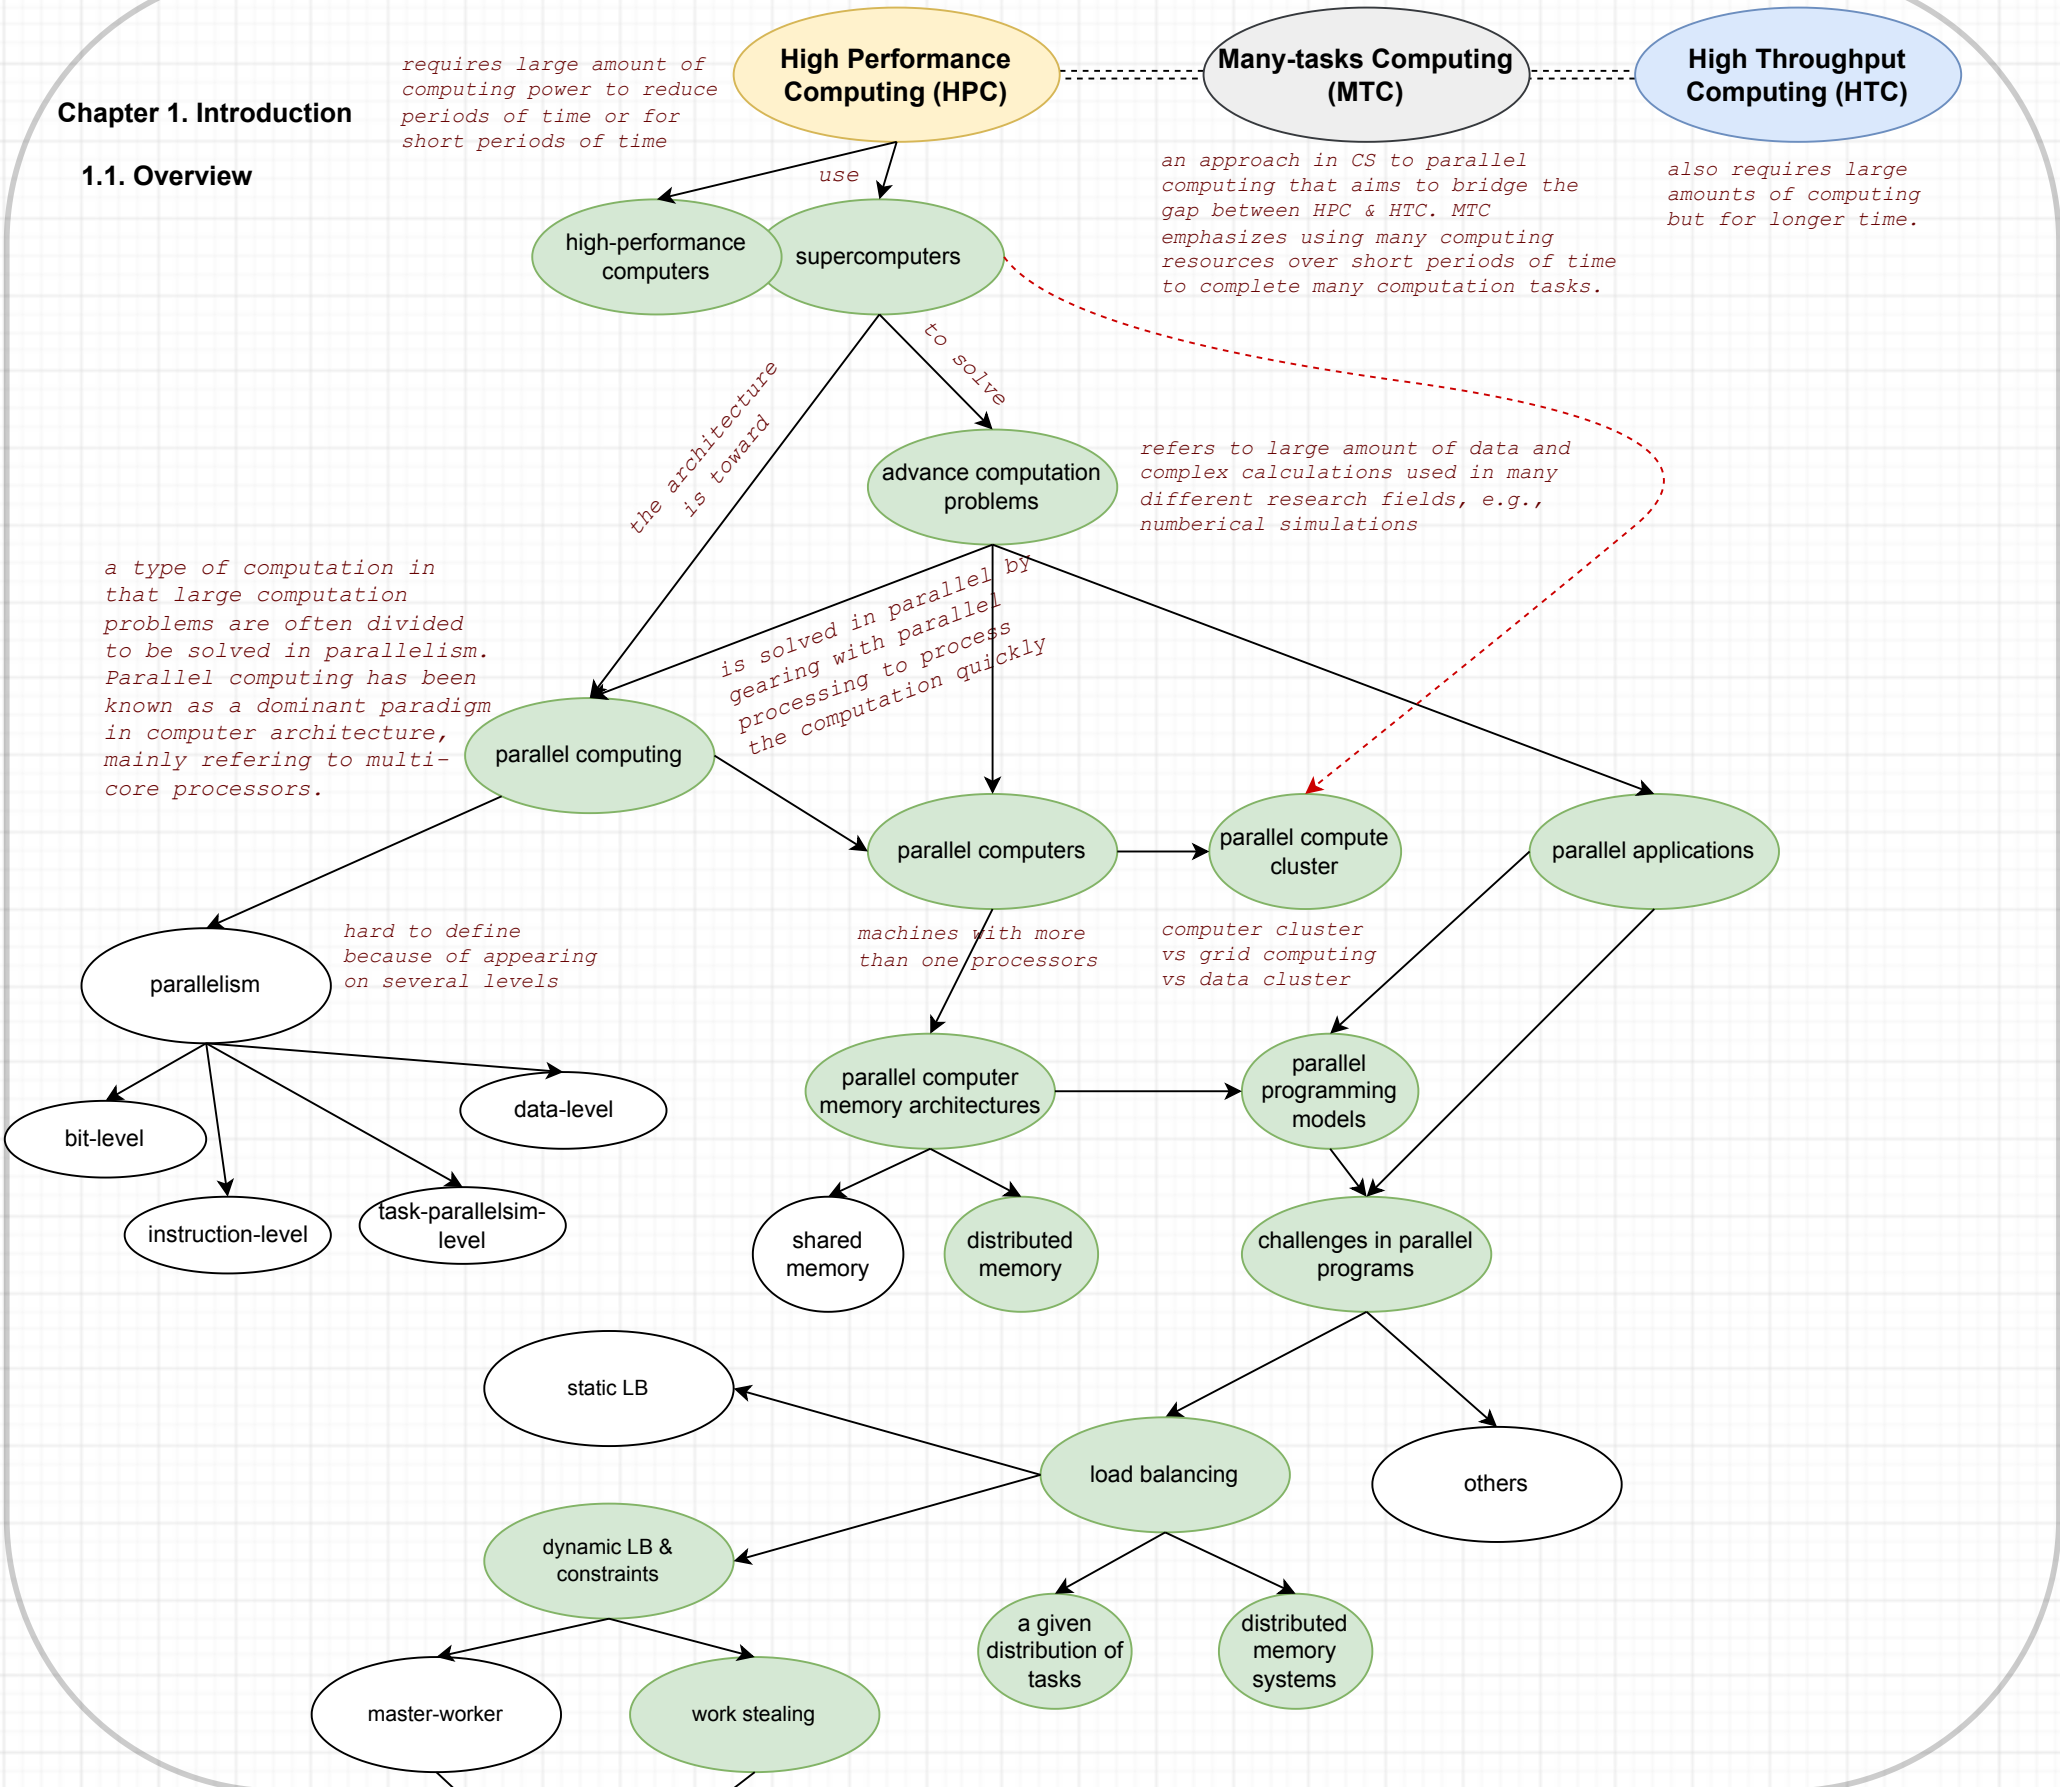
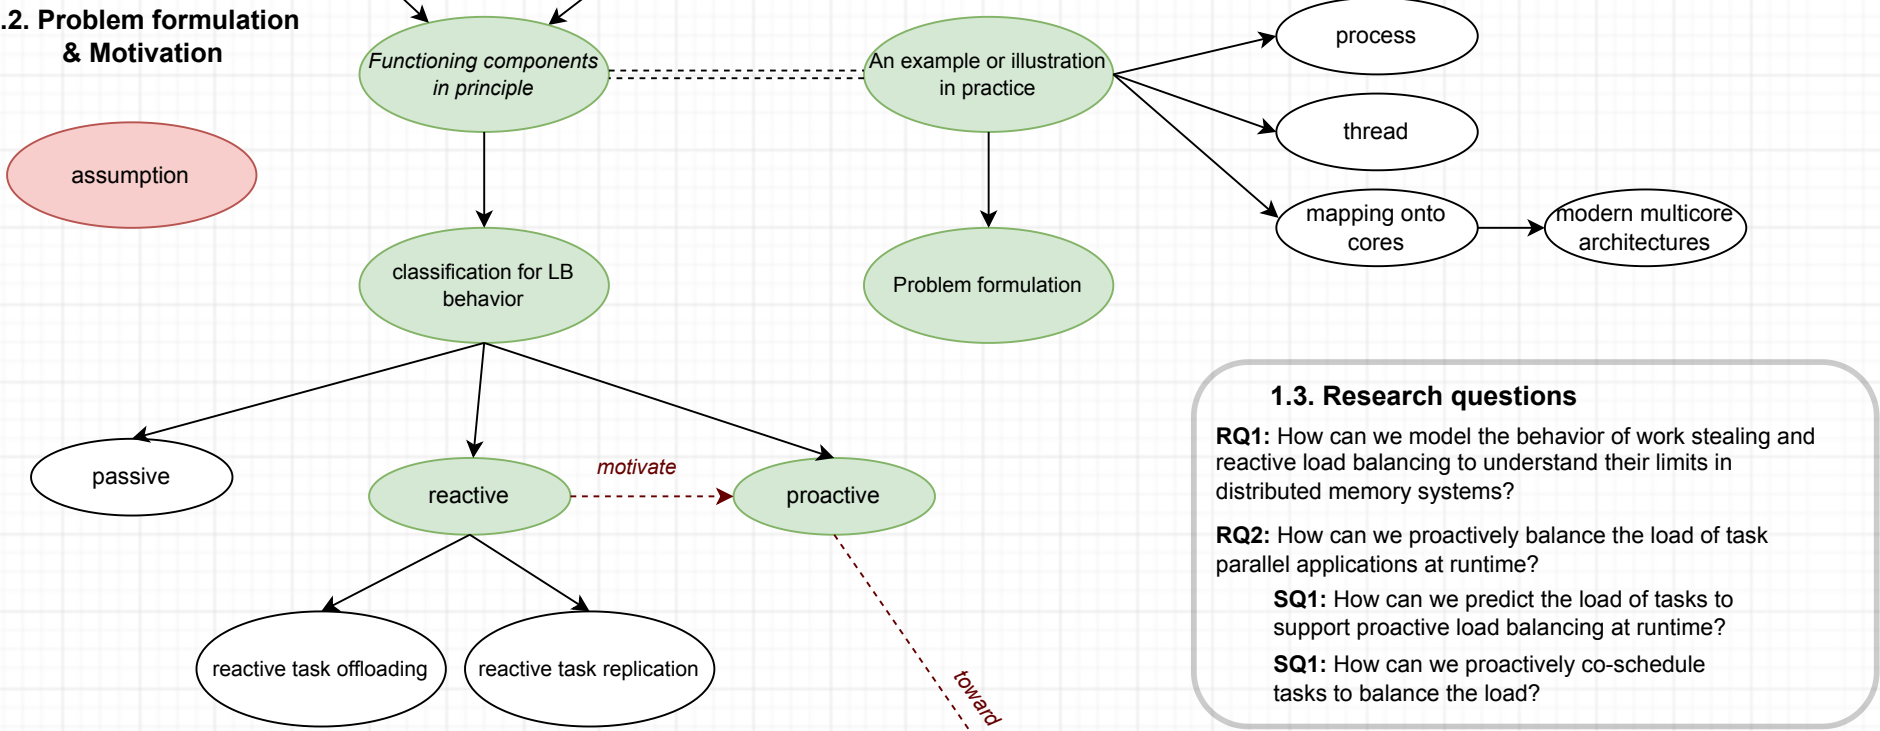


Chapter 1. Introduction

1.1. Overview



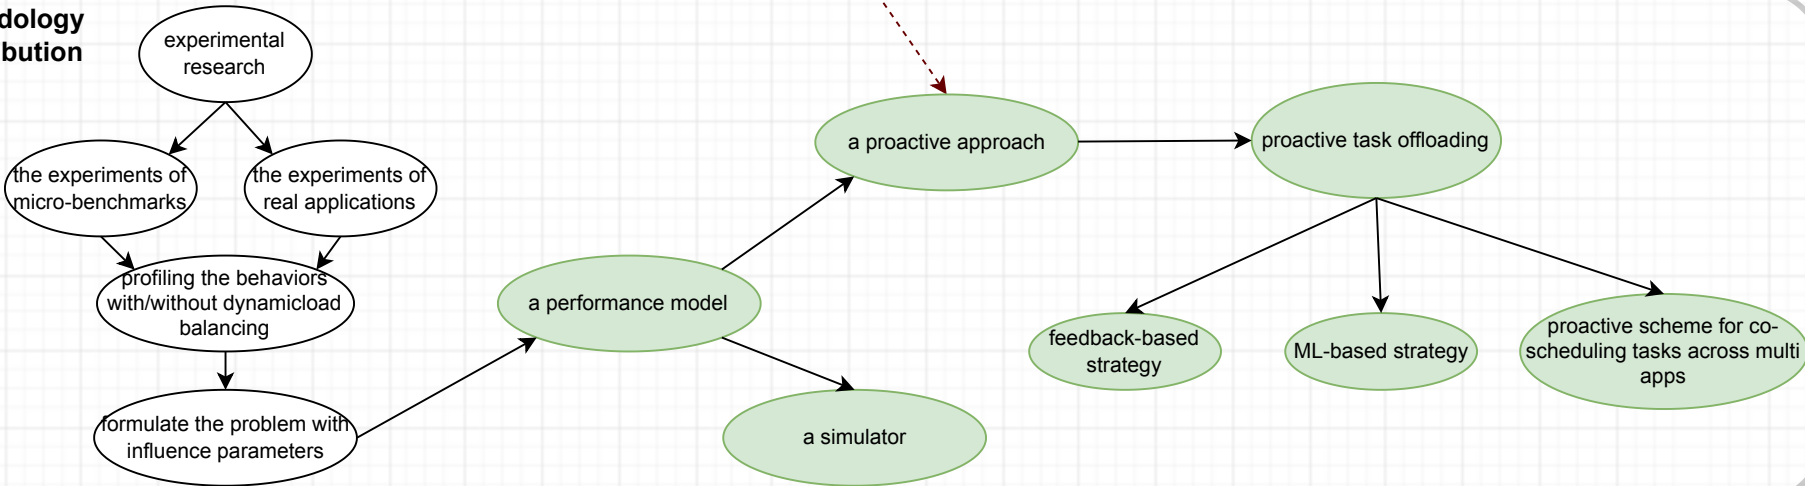
1.2. Problem formulation & Motivation



1.3. Research questions

- RQ1:** How can we model the behavior of work stealing and reactive load balancing to understand their limits in distributed memory systems?
- RQ2:** How can we proactively balance the load of task parallel applications at runtime?
- SQ1:** How can we predict the load of tasks to support proactive load balancing at runtime?
- SQ1:** How can we proactively co-schedule tasks to balance the load?

1.4. Methodology and Contribution



1.5. Publications

1.6. Thesis outline